

ABSTRACT

PURI, SHANIL . History Reuse Architecture: Reuse of Historical Computations across Iterations of a Converging Algorithm. (Under the direction of Xipeng Shen .)

In this data explosion era, there are myriad programs, which are executed repeatedly on large sets of data, consuming high amounts of energy and resources. Critical process reused on different data sets will no doubt help reduce the time and computations.

We study how Computation Reuse can be implemented for augmenting the performance of various classical algorithms such as K-means (Lloyd's Algorithm) and SGD Based SVM algorithm.

The classic K-means algorithm (Lloyd's algorithm) consists of two steps. For an input of 'n' data points of 'd' dimensions and 'k' initial cluster centers, the assignment step assigns each point to its closest cluster, and the update step updates each of the k cluster centers with the centroid of the points assigned to that cluster. The algorithm repeats until all the cluster centers remain unchanged in an iteration.

Another popular Algorithm of the same family is the Stochastic Gradient Decent based training for Support Vector Machine algorithm.

Our work with introduce a general architecture for comparing two or more data sets and get a probabilistic measure of similarity between them. We reason that if two data sets are similar, computation reuse from previous iteration of an algorithm for the similar data set will provide a good speedup in the execution of the algorithm for the current data set.

We first define the architecture to compute the above mentioned metric of similarity and then proceed to validate our above mentioned stipulation for the K-Means and the SGD based SVM algorithms.

In these explorations, we found challenges in 4 aspects: **data features**, **similarity definition**, **scalability**, and **reuse**. Data features, is the description of data- sets. Similarity definition (distance between two data-sets) is the way we would like to describe similarity between two data sets. Data feature and distance definition are used together for reuse instance selection from a program specific database. The selection of history record is the most essential part of computation reuse.

Since the method used to calculate distances between two data set's descriptions also should be defined based on the data features, the problem becomes even more complex. Scalability issues are related with number of instances in database, size of the target data set, and dimension of data-set. Goal of computation reuse is to save computation time and energy. In order to select a suitable history record, some extra computation is inevitably introduced. The dilemma is the trade off between the amount of computation reuse and the introduced overhead.

Reuse, thus is the process to decide what history information the program should reuse and how to reuse the pre-computed information.

© Copyright 2016 by Shanil Puri

All Rights Reserved

DEDICATION

To my parents.

BIOGRAPHY

The author was born in a small town ...

ACKNOWLEDGEMENTS

I would like to express my gratitude to my advisor Dr Xipeng Shen for the useful comments, remarks and engagement through the learning process of this master thesis. His constant guidance through the process invaluable, be it teaching me the correct methodology or discussing with me solutions for problems. His constant guidance was imperative in the successful completion of this thesis.

I would also like to take this opportunity to thank Dr Frank Meuller for providing me help with all the hardware that I needed for experimentation and also agreeing to being on my thesis committee.

I would also like to thank Dr. Min Chi for taking out the time for being on my thesis committee.

I would like to thank my loved ones, who have supported me throughout entire process, both by keeping me harmonious and helping me putting pieces together. I will be grateful forever for your love.

CONTENTS

List of Tables	vi
List of Figures	vii
Chapter 1 INTRODUCTION	1
Chapter 2 BACKGROUND and MOTIVATION	3
2.1 Benefits	3
2.2 Motivation	4
2.3 Background	4
2.4 Important Terminology	6
Chapter 3 Challenges	7
3.1 Data Feature Abstraction	7
3.2 Similarity Computation	8
3.2.1 Non Uniform Scale	8
3.2.2 Different Feature Sets: Feature Set Abstraction	8
3.2.3 Rotated Data	8
3.3 Efficient and effective comparison	9
3.4 How to use the historical datasets	9
Chapter 4 History Reuse Architecture	10
4.1 Overview	10
4.1.1 Feature Set Reduction And Normalization	10
4.1.2 Similarity Metric Calculation	11
4.1.3 History Storage Architecture	11
4.1.4 Matcher and Selector	12
4.2 Algorithm- HRu Architecture	12
4.2.1 Principal Component Analysis (PCA)	13
4.2.2 Histogram Generation	14
4.2.3 Similarity Metric Computation	15
4.2.4 Training Algorithm	15
4.2.5 Computing Best Match for Current Data-set (Run Time)	16
4.3 Discussion	17
4.3.1 Computation on Real Data and Incorrect Screening algorithm:	17
4.3.2 Custom CPP Implementation for Welch's Test	19
4.3.3 Random Sampling for Order comparison.	20
4.3.4 Use Of PCA Data for Welch's Test and PCA data for centroid computation in Training Run:	20
Chapter 5 Evaluations	22
5.1 Methodology	22
5.2 Experiments	22

Chapter 6	Related and Future Work	32
Chapter 7	BACKGROUND and MOTIVATION	33
7.1	Benefits	33
7.2	Motivation	34
7.3	Background	34
7.4	Important Terminology	36
BIBLIOGRAPHY		37
APPENDIX		38
Appendix A	LOREM IPSUM	39
A.1	A First Section	39
A.2	A Second Section	42

LIST OF TABLES

Table 5.1	Selection Hit Rate K-Means History Reuse	25
Table 5.2	Runtime comparison: Kmeans, Kmeans++ and Kmeans-HRU (<i>*All times in μS</i>)	28
Table 5.3	Median Iteration count comparison for Kmeans, Kmeans++ and Kmeans with History Reuse.	29
Table A.1	A table in the appendix.	42

LIST OF FIGURES

Figure 4.1	PCA Projection Example	13
Figure 4.2	Feature Set Abstraction	14
Figure 5.1	Hit-Rate (US Gas Sensor for $k=240$)	23
Figure 5.2	Hit-Rate (Caltec 101 for $k=600$)	24
Figure 5.3	Median Speedup, Small Data ($k=240$)	27
Figure 5.4	Median Speedup, Large Data ($k=600$)	27
Figure 5.5	Median Speedup, Small Data ($k=240$)	30
Figure 5.6	Median Speedup, Large Data ($k=600$)	31
Figure A.1	A figure in the appendix.	41

CHAPTER

1

INTRODUCTION

In this data explosion era, there are myriad programs which are executed repeatedly on large sets of data, consuming high amounts of energy and resources. Critical process reused on different data-sets will no doubt help reduce the time and computations. In this paper we will show introduce a new probabilistic model of measuring data similarities. Based on these data similarities we will show how critical computations may be reused across data-sets, saving energy and improving performance.

In this study we introduce the concept of Historical Computation reuse. This is an Intuitive Idea whose exploration is only spotted at best through history.

Computation Reuse is the process of deciding what history information the program should reuse and how best to reuse the pre-computed information. While the idea is simple and intuitive in nature, it promises big gains if correctly implemented with a wide variety of uses.

In this study we give an empirical study, showing that effectively computation reuse could enhance program performance.

Although, the idea of computation reuse is simple, there are many difficulties need to be solved to achieve efficient and effective reuse. In our explorations, we found out challenges in 4 aspects: data feature abstraction, similarity definition, scalability, and reuse. *Data features*, is the description of data- sets. Our algorithm thus must preserve the integrity of what the data represents. *Similarity*

definition (distance between two data-sets) is the way we would like to describe similarity between two data sets.

We used *Data feature abstraction* and *Similarity definition* together for reuse instance selection from a program specific database. The selection of history record is the most essential part of computation reuse because different record will lead to dramatically different computation amount. Because the most important properties of the input data may vary on different programs, it is difficult to decide a universal data feature and data distance definition. Since the method used to calculate distances between two data set's descriptions also should be defined based on the data features, the problem becomes even more complicated.

Scalability issues are related with number of instances in database, size of the target data set, and dimension of data-set. Goal of computation reuse is to save computation time and energy. In order to select a suitable history record, some extra computation is inevitably introduced into the computation. The dilemma is the trade off between the amount of computation reuse and the introduced overhead. For example, increase the number of instances in database will increase the probability of finding a good history record. However, it will also increase the introduced selection overhead. The size of target data set and the dimensionality of would also produce similar issues.

Reuse is the process to decide what history information the program should reuse and how to reuse the pre-computed information. Challenges in this field are mainly caused by the differences among programs and algorithms. For a specific program or algorithm, it might be a trivial solution while for another selecting historical data may be a complex problem. Our approach of building a global framework to work as a plug and play device makes the process even more complicated.

In this work, we investigate multiple solutions to address each of the challenges, and come up a framework, which we will apply to two of commonly used algorithms, namely: ***K-Means(Lloyd's Clustering algorithm)*** and the ***Stochastic Gradient Decent based Support Vector Machine*** for our explorations and providing proof of effectiveness of History Reuse and to validate the efficiency and the effectiveness of our algorithm.

The paper is organized in the following sections: *Chapter 2* gives background and definitions of terms we use in this paper. *Chapter 3* gives will give the motivation behind the work while *Chapter 4* will formalizes the framework and *Chapter 5* presents the challenges and solutions we explored in details. We present the experiment results in *Chapter 6*. Related work is discussed in *Chapter 7*.

CHAPTER

2

BACKGROUND AND MOTIVATION

2.1 Benefits

Computation reuse across executions benefits programs with long computation time most, especially converging algorithms that require multiple iterations to compute the desired results. Such a class of programs is essentially NP hard problems, which have no generic polynomial time algorithms. Another benefit may be improved accuracy. Converging algorithms belong in the NP class of problems and do not have optimal solution. Through appropriate history result reuse, numbers of iterations of computations could be saved while improving the accuracy of the algorithms. The key point of computation reuse on different data is finding suitable history information to reuse. Thus the above problem stated problem essentially boils down to introducing a probabilistic method of calculating data-set similarities quickly and accurately. With databases containing a number of instances, data- set feature is a key component of each instance, and history computation result of the program on this data set is the value. Then, through computing distance from current data set to each instance, our framework could select the instance, which has the highest probability to provide most computation reuse.

2.2 Motivation

Optimization of converging Algorithms such as K-means clustering and SGD based SVM are highly beneficial. Algorithms of this category are used frequently and have a multitude of applications in real world machine learning and big data processing. In this age of data explosion this, thus is a highly lucrative area of research. Some of the motivations for this work may be listed as follows:

- Frequently Used: These algorithms are used frequently to tackle real world problems and thus even small improvements can have a significant impact in performance improvement.
- These algorithms also have wide areas of real world application ranging from machine learning to big data problems.
- These algorithms are ideal suited for such optimizations, as their speed of convergence and accuracy is directly dependent on the starting points for the algorithm and thus can be used as proof for the benefits of history reuse easily.

2.3 Background

For most part of the generation of our framework for Historical Dataset selection we use the classical *K-Means algorithm (Lloyd's Algorithm)* for the finalization and testing of our *History Reuse architecture*. We then test the architecture with the *Stochastic Gradient Decent based Simple Vector Machines* to prove the global viability of our algorithm.

The classic K-Means algorithm (Lloyd's algorithm) consists of two steps. For an input of 'n' data points of 'd' dimensions and 'k' initial cluster centers, the assignment step assigns each point to its closest cluster, and the update step updates each of the k cluster centers with the centroid of the points assigned to that cluster. The algorithm repeats until all the cluster centers remain unchanged in an iteration. Because of its simplicity.

Because of its simplicity and general applicability, the algorithm is one of the most widely used clustering algorithms in practice, and is identified as one of the top 10 data mining algorithms (Wu et al., 2008). However, when n, k, or d is large, the algorithm runs slow due to its linear dependence on n, k, and d. There have been a number of efforts trying to improve its speed. Some try to come up with better initial centers (e.g. K-Means++ (Arthur and Vassilvitskii, 2007; Bahmani et al., 2012)) or parallel implementations (Zhao et al., 2009). This thesis will look to present more on this approach by exploring an avenue not much pursued before, namely historical data set cluster center reuse for K-Means Initialization.

Prior efforts in this direction include: K-Means++(Arthur and Vassilvitskii, 2007; Bahmani et al., 2012), K-Means Initialization Methods for Improving Clustering by Simulated Annealing (Gabriela Trazzi Perim et al. 2008), An optimized initialization center K-Means clustering algorithm based on density (Xiaofeng Zhou et al. 2015). These prior methods, while having made a significant contribution, have failed to replace the Lloyd's algorithm which still remains the dominant choice in practice exemplified by the implementations in popular libraries, such as GraphLab (Low et al.), OpenCV, ml-pack (Curtin et al. 2013) and so on.

For an alternative algorithm to get widely accepted, we believe that it needs to meet several requirements: (1) It must inherit the level of trust that Lloyd's algorithm has attained through the many decades of practical use; (2) it must produce significant speedups consistently; (3) it must be simple to develop and deploy. The previous methods of initializations though have an innate flaw. The initialization points are chosen at random, albeit sometimes with certain steps to increase probability of getting good starting points (e.g.K-Means++ (Arthur & Vassilvitskii, 2007; Bahmani et al., 2012)). Thus the probability of finding the best starting point is upper bound by the error rate of at least the first point chosen at random. This, thus restricts the best choice of points for our current algorithm. Historical dataset cluster center reuse atleast theoretically can initialize the data set which optimal or close to optimal starting centers if an exact or close to exact match for the dataset is found.

The previous proposals unfortunately fall short in at least one of the three requirements. The approximation methods (Czumaj & Sohler, 2007; Sculley,2010; Philbin et al.,2007;Guha et al.,1998; Wang et al.,2012), for instance, produce clustering results different from the results of the standard K-Means. It is possible that their outputs could be good enough (or even better) for some usage, but in general such a confidence is yet to be established, in both theory and practice: Users still face the uncertainty on whether the output from these algorithms is good enough on an arbitrary dataset for real usage, for which Lloyd's algorithm has been practically proven to work correctly.

This thesis will introduce formalize the History Reuse Architecture. We will then use this History Reuse to introduce a new Initialization methodology for the K-means algorithm: "Historical dataset center reuse for K means initialization", an enhanced K means implementation which aims to optimize the K means algorithm by aiming to choose the best possible starting points for the K-Means(Lloyd's algorithm) for faster convergence. Since the only modification that are being made are in the initialization step of the algorithm it stands to reason that the algorithm would continue to uphold the same standards as the standard K-Means algorithm. We will also use our History Reuse Architecture and test it for SGD based SVM to prove the global application of our generic architecture.

2.4 Important Terminology

The following are Important terms for this work and are used through out the later chapters.:

- **Data Source:** These are the actual sources of Data Set repositories from which we source our data for testing purposes.
- **Data Set:** Data on which the actual algorithm is run after dividing the data from the *Data Source*. Each *Data Set* is built up of multiple *Data Items* or *Data Points*.
 - *Current Data Set:* Represents the data set on which the current iteration of the algorithm is to be run.
 - *Historic Data Set:* Data Set present in the History Data Base for which results have been computed in previous iterations of the algorithm, making the data set one of the viable candidates for *History Reuse* for *Current Data Set*.
 - *Data Item / Data Point:* Single row in Data set. Represents a single data point in the larger *Data Set*.
 - *Dimentionality(d):* Number of columns used to represent a single *Data Item*.
- **History Data Base:** Data base for storing all data sets that may be candidates for *History Reuse* for *current data set* and for which final results have been calculated in previous iterations of the algorithm.
- **Cluster Count (k):** The total number of cluster in which the *Data Set* is to be clustered when using the *K-Means Algorithm*.
- **SGD based SVM:** Stochastic Gradient Decent Based Simple Vector Machine.

CHAPTER

3

CHALLENGES

History Reuse seems an intuitive solution to many problems. While the concept in itself is simple enough: reuse some computations for previous runs for current run of an algorithm, the implementation of the same provides quite a challenge. Some of the major challenges faced are: Data Feature Abstraction, Similarity definition, Scalability, and Reuse.

3.1 Data Feature Abstraction

Each *Data Set* is categorized by a set of *features* in the form of columns, assuming the data set is represented as a 2D array. In such a case not all features for the data set hold equal importance in the data set categorization. Thus the first challenge faced by us was the ensuring that we use the most important features only for computation of our *History Reuse Data Set*. Use of too many features for the computation of our Historic Data Sets may effect efficiency, while on the flip side, use of too few data set features may result in the loss of the meaning of the data set itself thus invalidating its candidacy for History Reuse. This thus was one of the major challenges face by us in the computation of Historical Data Sets for Computation reuse.

3.2 Similarity Computation

Another important feature requirement for our problem statement was to come up with a uniform metric for feature set comparison. For this we propose a “Probability based” metric for similarity between data sets. By this metric we can make a quick yet accurate assessment regarding the degree of similarity in between data sets. Obviously the best choice of historical data set would be the one with the highest probability of being similar to the current data set. Defining such a metric though poses its own challenges, namely: data sets used for comparison may have different scale (may not be normalized), may not have the same important features or may not be along the same axis of projection. These and more issues make the definition of a Similarity metric a difficult task. Failure to solve all the challenges mentioned above would lead to the failure of our *probabilistic similarity metric*. Following sub sections will shed a little more light on the challenges in :

3.2.1 Non Uniform Scale

Data sets present in the *History Data Base* and the current data set may not have the same scale, i.e distributions and variances. This essentially means, any similarity computations between the two data sets would essentially be meaningless as computation reuse across such data sets may be impossible. This thus presents the first challenge of normalizing the data sets on to the same plane to provide a common platform for similarity computation.

3.2.2 Different Feature Sets: Feature Set Abstraction

Another major challenge faced is that the most influential features present in *historical data sets* and the current data set may be very different. Since the class of algorithms targetted by our algorithm are often highly dependent upon the feature sets for final results, the matching of data sets based on only the most important features becomes imperative for good reuse results. This, thus presents the challenge of analyzing both the current data set and the historical data set for the extraction of the most important features. This must be done in real time and must be both efficient and effective. as mentioned in *section 1*, we must also ensure the optimal use of the features to ensure meaning full comparisons.

3.2.3 Rotated Data

Often the axis of projection for current and historical data may often in different planes, and while the distributions may be similar for data, their being on different planes altogether makes similarity computation challenging. Thus one of the major challenges in computing similarity was to ensure

both *historic and current data set* be in the same plane. This again must be done at run time, so as to ensure that both the historic and current data and the historic data are on the same plane (*plane* for current data is known only at run time.) This further poses the challenge of an efficient method for the planar normalization of the data sets being compared.

3.3 Efficient and effective comparison

The above challenges seen clearly pose the challenge of efficiency and effectivity. The loss of either of the two would essentially entail the failure of any proposed algorithm.

Efficiency is desired since a lot of the above challenges must be solved in real time as they require analysis of the data set used in the current iteration of the algorithm.

Effective is desired since a non effective data set may provide a bad historical data set for computation reuse. Our experimentation has shown us, use of badly matched historical data sets tend to cause severe punishment in terms of both run time and quality of results for our tested algorithms.

3.4 How to use the historical datasets

The last challenge faced is how to reuse computation from the selected history data set itself. For example in case of the K-Means algorithm we get both the labels, as well as cluster centroids for historical data sets. The use of labels though quicker in initialization may produce different results depending upon the order of the data points in the data set. Another example maybe, the data item count for historical data set may be different from the current data set and thus using labels directly may simple cause the failure of the algorithm to divide the data into the required number of clusters. On the flip side too few data items in historic data set may cause some data items in current data set to not being assigned any cluster at all. The reuse of computed centroids to regenerate labels on the other hand, while slower, will always ensure the correct label initialization irrespective of the order of the data points in the data set, or the count of the required data labels etc. Thus we can see that the correct use of Historical data is an equally challenging problem.

CHAPTER

4

HISTORY REUSE ARCHITECTURE

4.1 Overview

For any framework to be successful globally in the computation of History Reuse needed to effectively and comprehensively solve all the challenges mentioned above. In addition it also needed to be directly translatable and not be over dependent upon the algorithm in question. Keeping the above framework requirements in mind we finally settle on the framework components as follows:

4.1.1 Feature Set Reduction And Normalization

The aim of this section of the framework is to compress the data to the highest extent possible while maintain the meaning of the data set. For this purpose we decided to take a leaf out of the Image processing/statistics books. This step consists of two major parts:

- **Dimension Reduction:** In this part of the algorithm we essentially reduce the total dimensions of the data set to the minimal possible without losing the meaning of the data set. This is achieved by the use of Principal Component Analysis (PCA). PCA is a method, which takes in a data set, and then proceeds to return a modified data set such that the dimensions in the updated normalized data set are arranged in descending order of variance. This way we can

take the top few dimensions for our computations.

- **Point Count Reduction:** In this part of the algorithm we essentially aim at selecting the optimal sample set of data-points from the data set for our computations. This is essentially achieved by dividing all the points into buckets. We then proceed to pick the top-k highest populated buckets to get the highest density range of the data set, while at the same time greatly reducing the total no of data-points that would be needed to be computed in our data-set.

4.1.2 Similarity Metric Calculation

- Once we have evaluated the reduced data-set using steps mentioned in section 5.1, we proceed to calculate the probability based similarity metric. We achieve this using **Welch's Test** for non-parameterized data for null hypothesis testing.
- Our assumed null hypothesis for any compared data sets is that both data sets are exactly similar to each other. The Welch test thus gives us a probability metric that states that the difference in data sets is due to chance. Thus higher the probability of the difference in data sets being up to chance, the better is the probability that the two sets would be similar.
- We take individual dimensions from the two data-sets and run Welch Test on them to get a probabilistic measure of their similarity on a per dimension and a cumulative sum across dimensions. We then rank the data sets with respect to each other based on the computed probabilistic metrics. Now this is an important aspect of our computation because this is the algorithm that we used for our probabilistic metric calculation, which is in turn used for ranking all data sets relative to each other.

4.1.3 History Storage Architecture

Once we have computed the above mentioned values for our data-set, we store it in memory as objects. Each historical data objects holds its original data, PCA metadata (eigen values and eigen vectors etc), bucket-wise histogram and the reduced data set. In addition to this each object also stores a score of the probabilistic similarity it holds with each of the other historical data sets and maintains them in non-increasing order of probability. This is done for quick access of data sets for computations when we are comparing real time data with historical data-sets.

4.1.4 Matcher and Selector

This is by far the most time critical part of the framework. It needs to be quick because this is the function that is responsible for matching the current data set with the sets in the historical data-set and find the best match for computation reuse selection in real time. Since, historical data may be large, we need a quick way to find the closest match from the historical data set. Our framework goes about doing this in the following two passes:

- In the first pass we do the PCA computation for our current set. We then proceed to use the histogram made by the PCA data points for quick distance comparison. We take the highest populated top-k buckets and do a distance computation with points in similar buckets in other data sets. This gives us one candidate for History Reuse. Another candidate is calculated by computing distance as mentioned above but in this case instead of taking highest populated buckets from History Reuse candidates, we now use buckets having closest ranges to the selected buckets for current data set. This is done to estimate similarity in distribution. Doing this we now get another candidate for history reuse. This part of the algorithm runs linearly without much time delay. Thus we can afford to do this kind of matching with all the historical data sets and get the approximate closest match.
- We now use the above computed candidate historical data sets along with the current data set to compute the probabilistic metric of similarity between the data sets and then proceed to compute the same metric for the top three closest matches to the historical data sets (pre-computed and stored.) We now use the data set with the highest probabilistic similarity to select data set for computation reuse. This step gives us the final candidate for History Reuse.

4.2 Algorithm- HRu Architecture

The final algorithm maybe divided into two major categories: Training and Run Time.


- **Training Algorithm:** The training algorithm is used to prepare and store data so as to have highest availability of reuse components for all candidates. This part of the algorithm is carried out offline and thus does not effect the run time of the algorithm when run for most current data set.
- **Run Time Algorithm:** The run time algorithm is responsible for choosing the best match historical data set to be used for Computation Reuse. This part of the algorithm is online, thus it has a direct impact on the run time for the algorithm when run for current data set. This

part of the algorithm must be quick and must ensure that the computation time required for historical data set selection not overshoot the time for benefits gained by such history reuse.


Both the training algorithm and the run time algorithm use some complex methodologies to ensure the best possible approach for historical data set selection. These are defined in the following sub sections.

4.2.1 Principal Component Analysis (PCA)

PCA and project data set points on thus computed eigen vectors. PCA essentially extracts the top n most influential *features* of a data set. For our algorithm we choose the **top 3 principle components**. PCA returns eigen vectors for the three chosen principal axis (or components) for our data set based on maximum variance for each feature set (represented by a single column). We then proceed to project the data onto the new eigen plane using the above computed eigen vectors. This step reduces the dimensionality of our data into a fixed three dimension space. **figure 4.1a** shows an example data distribution in a 2-D plane and **figure 4.1b** shows the same data projected into a 2-D eigen plane. Since eigen planes are always coplanar, this thus normalizes the data along the same axes.




(a) Example: Real Data with 2 Dimensions




(b) Example: Data Projected on Eigen Plane

Figure 4.1 PCA Projection Example

We can also see from 4.1b that maximum variance of data is along the **pc1 axis**. This makes the **pc1 axis** the 1st principle component of the data. Using this property of PCA we can choose the most important *features* only from a data set while excluding the less important *features*. We see from **figure 4.2a**, real data needs both the **x and y plane** to represent the features (spread) of the data. On the other hand in **figure 4.2b** we can see that most of the variance for the data items has been condensed with negligible variance along the **pc2 axis**. All features of the data can now be abstracted onto the **pc1** plane. This method is the method we use for data abstraction while maintaining the features of the data set.



(a) Example: Variance for Real Data on X and Y axis individually



(b) Example: Variance of Data on PCA axis individually

Figure 4.2 Feature Set Abstraction

4.2.2 Histogram Generation

Once the dimensionality of the data has been reduced using PCA, we need to reduce the point count ensuring we use the most dense distribution areas for data set comparisons. We tackle this problem by the use of Histograms. Histogram is then generate on a per dimension (*di*) basis with each dimension having 16 or 32 buckets with ranges computed as follows :

$$bucket_count = 16 \text{ or } 32$$

$$dim_bucket_range_{di} = max_{di} - min_{di} : \forall di \text{ in } (1 .. 3)$$

$$bucket_size_{di} = dim_bucket_range_{di} / bucket_count : \forall di \text{ in } (1 .. 3)$$

$$bucket_range_j = [(bucket_size * bucket_j + min_{di}), (bucket_size * (bucket_j + 1) + min_{di})]; \\ \forall j \text{ in } (1 .. n)$$

(4.1)

This will give us the buckets with their ranges for all dimensions. We then proceed to generate *three* histograms. Each data point is put into a bucket based on its value for that particular dimension.

Eg: let *three points* be represented in a 3-D eigen plane by coordinates as follows:

$$p1 = (1, 2, 3) \quad p2 = (-1, 0, 1) \quad p3 = (2, 3, 4)$$

Now, let there exist 2 buckets per dimension with Ranges as follows:

$$b_{11} = [-1, 2) \quad b_{12} = [2, 4) \quad b_{21} = [0, 2) \quad b_{22} = [2, 4) \quad b_{31} = [1, 3) \quad b_{32} = [3, 5)$$

Using the above points and bucket ranges we will now generate **3 Histograms**(one per dimension) with **2 buckets per histogram**. The histograms may be represented as follows:

$$h_points_1 = b_{11} : (p1, p2); b_{12} : (p3); \text{Histogram for first dimension}$$

$$h_points_2 = b_{21} : (p2); b_{22} : (p1, p3); \text{Histogram for second dimension}$$

$$h_points_3 = b_{31} : (p2); b_{32} : (p1, p3); \text{Histogram for third dimension}$$

Thus we can see from above example how each item becomes a part of a bucket if its coordinate for the given dimension lies in the bucket range for that dimension.

4.2.3 Similarity Metric Computation

We now use the Welch's Test to compute the probabilistic metric for similarity and to rank all history data sets wrt to each other and store in decreasing order of similarity. (Note: Will be used in run time historical reuse data set computation.) Welch test is computed for each dimension of the two sets being matched and the summation of the score for all three dimensions is as Similarity score for the two data sets. Data sets are then ranked with each other based on the similarity score. Higher score means a better and match and vice versa.

4.2.4 Training Algorithm

This part of the algorithm is used to prepare and store data so as to have highest availability of reuse components for all candidates. We try and solve the challenges mentioned in chapter 3 using the methods mentioned in 4.1. Let our history database consist of "n" data sets named: HD_1, HD_2, \dots, HD_n . The following Steps are done for each HD_i for (i in 1 to n) :

- **PCA:** Compute eigen vectors for HD_i and project data on the eigen vectors to normalize in eigen plane and reduce dimensionality as shown in 4.2.1. We reduce all data into 3-D space for comparison purposes.
- **Histogram Generation:** Categorize PCA data for HD_i buckets for histograms as shown in 4.2.2. One histogram is generated per dimension of data, thus we have a total of 3 histograms for each data set. Each histogram comprises of a total of 32 buckets.
- **Relative Ranking:** Use Welch Test to rank all history data sets wrt to each other and store

in decreasing order of similarity. (Note: Will be used in run time historical reuse data set computation.)

- **Reuse Data for Data Set:** Run Algorithm of choice for Data Set and store *Computation Reuse Data*. Eg: In case of the K-means algorithm we choose the *final centroids* computed by the K-Means Algorithm for current data set.

4.2.5 Computing Best Match for Current Data-set (Run Time)

This is the part of the algorithm

- **PCA:** Compute PCA for current Data Set similar to 4.2.4.
- **Histogram Generation:** We again characteris PCA data into 3 histograms with 32 buckets each as done in 4.2.4.
- **Matcher and Selector :**The best match data set selection process is divided into 2 parts: screening and final selection. Screening is used to prune our Historical Data Sets to a smaller subset from which we can select our best match Data Set for computation reuse using the final selection step.

The two parts of the Matcher and Selector part of our algorithm may be defined as follows:

- **Screening:** This is the process where we select our initial candidates for the next step of our selection algorithm. This part of the algorithm essentially estimates the similarity in distribution for the current and corresponding historical data sets using the Histograms generated in the previous steps. Variance Similarity is estimated as follows:
 1. Take top-k most populated buckets for current data set and select corresponding top-k bins from all history data sets and top-k buckets with closest min and max to current data-set top-k bins.
 2. Find ED for these data points between current data set and all historical data sets.
 3. Choose Historical Data Sets with minimal distance as initial “Best Match” for both top-k buckets by rank and top-k buckets by range.

Pseudo Code for an understanding of how our screening algorithm runs can be seen in 1.

- **Find Best Match (Similarity Metric => Probabilistic Score):** One we have found the best match candidates from out screening step, we simple calculate the Welch Test score for each dimation (dim_i) for each pair of current data set and chosen historic data sets.

1. Find “Similarity Metric” as defined in 4.1.2 between current data set and above chosen “best match” data sets using Welch’s Test for all dimensions of all three data sets.
 2. Similarly find “Similarity Metric” for top two relatively ranked data sets for current best match and top one for second best match.
 3. Choose Data Set with highest Probability Score.
- **Reuse Computations:** In this section we use the computational data we had stored for history reuse in the training run for the historical data set selected. For instance, in the case of the K-Means algorithm, we now use the centroids from the chosen historical data set that were computed during the training run for historical data.

4.3 Discussion

Our final algorithm was reached at after a fair few trials and errors. Some of the major approaches used apart the latest approach for major challenges may be defined as below:

4.3.1 Computation on Real Data and Incorrect Screening algorithm:

4.3.1.1 Implementation

- In this method I had initially use bin wise reduction on initial (real) data and then used PCA for dimension reduction of these reduced data points for the calculation of eigen vectors and eigen values only.
- I had then proceeded to run Student’s T-Test for relative ranking in in training Run on real data.
- For Run Time computation I had used only the distance between the eigen vectors for initial screening, choosing the data set with smallest difference in eigen vectors as initial Best Match data set.
- I had then proceeded to Use Student T-Test on real data for computation of the probabilistic metric using all dimensions for the computation of the same.
- Python Numpy Libraries had been used for Student T-Test computation.
- History Reuse component Computations during training run was done on real data of historical data set instead of PCA data.

```

least_distancerank = MAX
least_distancerange = MAX
dim_distancerank = 0
dim_distancerange = 0
HD_Rankselected
HD_Rangeselected
forall the HDx in History Database do
  forall the Dimi where i in 1 .. 3 do
    forall the Bucketj in Top 3 Buckets(Dimi) in Decreasing Order of Item Count for
    current data set do
      forall the ptcd and ptHDi in points(Bucketj) do
        distancerank = dist(ptcd, ptHDi);
      end
      Bucketk = Bucket in HDi where range(bucketk) range(bucketj)
      forall the ptcd in points(Bucketj) and ptHDi in Bucketk do
        distancerange = dist(ptcd, ptHDi);
      end
      end
      dim_distancerank = dim_distancerank + distancerank
      dim_distancerange = dim_distancerange + distancerange
    end
    if dim_distancerank < least_distancerank then
      least_distancerank = dim_distancerank;
      HD_Rankselected = HDx;
    end
    if dim_distancerange < least_distancerange then
      least_distancerange = dim_distancerange;
      HD_Rangeselected = HDx;
    end
  end
end

```

Algorithm 1: Screening Algorithm**4.3.1.2 Reasons for failure**

- Data was not normalized thus computation would not be correct.
- Student T-Test worked only with Gaussian Distributions. Garbage value of non Gaussian Data.

- Eigen Vectors of two data sets may be orthogonal yet PDF (probability density function) may be close enough such as to generate similar clusters.
- Use of labels for direct initialization was flawed in the sense that if the data points were jumbled they would produce the wrong order of labels thus still providing a bad match.
- Bin Wise distribution of data points was an expensive operation and computation overhead increased with increase in dimensionality of data.
- Student T-Test had to compute for multiple dimensions of data and was a time expensive computation.

4.3.2 Custom CPP Implementation for Welch's Test

Some of the earlier seen issues were corrected in this section as I noticed that a lot of the run time improvement was overshadowed by the time taken for choosing the historical data set. I also noticed that the Student T-Test was not reliable for non gaussian data and failed completely in case of different data set sizes. This iteration was also mainly about ensuring quicker run time for the selection algorithm, so incremental updates were made to optimize the same. Some of the updates made may be enlisted as follows:

- Implemented Custom Cpp implementation for Welch's Test to over come overhead created by using python libraries for it and calling python script from Cpp.
- Removed the distribution of data into bins for data point reduction as it had a big overhead in computation and CPP Welch Test Libraries scaled well for larger data sets.
- Changed to use of Welch's Test as compared to Student T-Test as Welch's Test works with non Gaussian distributed data as well as compared to Student T-Test which makes assumptions of data distribution being Gaussian in nature.
- For Run Time computation I had used only the distance between the eigen vectors for initial screening, choosing the data set with smallest difference in eigen vectors as initial Best Match data set.
- Centroid Computation during training run was done on real data of historical data set instead of PCA data.
- Labels for Best Match historic data set were used as-is for current data set.

4.3.3 Random Sampling for Order comparison.

Experimentation on the approach showed non-reliable best match selection. Also I saw that often the best match may not even yield best results. One of the major reasons for this was the use of the incorrect use of historical data. For instance, incase of the K-Means algorithm the use of labels instead of the computed centroids lead to a dependence on the order of the items in the historic data set. While the two data sets may be similar and produce similar clusters, History Reuse with this method would still fail as initially the data points in current data set may get labeled incorrectly. I also realized that comparison of non normalized data was in its very essence. I tried to correct the above issues with the following methodology:

- Changed to use of PCA data for most computations.
- Projected Historical data set points on Current Data Set eigen Vectors. Used distance between data points of Current Data Set and historical data sets by randomly sampling 10% of data sets against each other. Chose Data Set with minimum distance. This was done to try to use Historical data set with most point order similarity (and thus genrated label order similarity) in conjunction with overall data set similarity.
- Welch's Test was still used across all dimensions of actual data Similarity Metric Computation.
- Labels were used as-is for real data.

This implementation though corrected some of the above mentioned issues, it in its turn generated new issues:

- Use of labels for direct initialization was flawed in the sense that if the data points were jumbled they would produce the wrong order of labels thus still providing a bad match.
- Random Sampling was a bad way to judge the order of the labels that would be generated by K-Means for data set.
- Eigen Vectors of two data sets may be orthogonal yet PDF may be close enough such as to generate similar clusters.

4.3.4 Use Of PCA Data for Welch's Test and PCA data for centroid computation in Training Run:

Experimentation still showed both the history reuse to be inconsistent and the overhead of computing historical data set was still quite large. i also realized that while Eigen Vectors of two data sets

may be orthogonal yet PDF may be close enough such as to generate similar clusters. To correct the above issues I used the following methods:

- Used PCA data in training run for Centroid computation of historical data.
- Used PCA data for Welch's Test based "Similarity Metric" computation.
- Still used labels from best match historical data set as initialization for current data set.
- Stopped projecting historical data sets data on current data set eigen vectors, instead used self projection which could be done offline.
- Used Random Sampling for initial estimation of best match data set.

The use of Labels was still flawed. Also the use of random sampling was not a very effective methodology for screening as it would often lead to the selection of bad candidates.

CHAPTER

5

EVALUATIONS

5.1 Methodology

5.2 Experiments

To demonstrate the efficacy and efficiency of our program we evaluate our approach by testing it on various large real world data-sets and compare our algorithm to two most frequently used and vastly accepted K-Means initialization algorithms: standard K-Means (Random Select initialization also known as Lloyd's Algorithm) and K-Means++ initialization (Weighed Probability based initialization based on squared of distance from initial randomly selected center). Both the above algorithms are implemented in the OpenCV library, which has been used as the standard library for Lloyd's Algorithm implementation. We run all three algorithms on the same data-sets with the same error rate for convergence.

We have also compared the historical data-set selected by our algorithm to all other data-sets available in history to show the accuracy of our algorithm in selecting the best available data-set.

Consistent Selection of Best Match Historical Data-set from History: The experiments run, show that our algorithm is able to consistently able to select one of the top 5 data-sets for history reuse from available historical data-sets. By consistent, we mean that our algorithm is able to select




Figure 5.1 Hit-Rate (US Gas Sensor for $k=240$)

one of the top 5 available data-sets with hit rate of above seventy percent across data-sets irrespective of the size and the cluster count. This is expected as we aim to compare data set based on similarity while the cluster count is not taken into consideration in the matching process. **Table 5.1** shows the hit rate for our algorithm. We can clearly see from the table that our algorithm is able to select one of the top 5 best data-sets for historical reuse from the available data-sets consistently irrespective of the cluster count and the total historically available data-sets. This thus validates our theory for probabilistic selection for matching data-sets using the **Welch's Test** for null hypothesis testing for similarity of data-sets. We are also able to see clearly that our selection criteria continues to perform well across a wide range of cluster counts and thus validates our hypothesis that data-set similarity should be used as the primary measure to gauge quality of data-set for history reuse without the explicit need for taking into account the total number of clusters (k) the data-set is to be clustered into. We have chosen to limit the total number of clusters to 256 so as to have meaningful clusters for data-sets of relatively small sizes. Columns 3, 4 and 5 show us how well our data-set selection algorithm works across various cluster indexes and also how little variation is seen in the performance of the data-sets selected by our algorithm irrespective of the number of clusters (k). We also show that our hit rate is a minimum of 70 percent (Some error is to be expected as the selection




Figure 5.2 Hit-Rate (Caltec 101 for $k=600$)

algorithm works on probabilistic model for making best guess.)

Figures 5.1 and 5.2 show how our selected historical data set performs in comparison to all other data sets available for selection in our history data base. We see from them, the accuracy of algorithm in consistently choosing one of the best match data sets from history with high accuracy ($> 75\%$). We see that our algorithm chooses either the best match data set or data set close to best match. This proves the accuracy of our algorithm.

Consistent Scalable approach to Selection: Our experiments also show that our method scales well with the increase in dimensionality of data (d), the size of the data-set(n) and the cluster count(k). **Table ??** shows our experimental results for various data-sets with varying dimensionality and sizes. The below table clearly shows that for any given data set the selection time is completely independent of the cluster count (k). This is in complete contrast to both the *random K-Means (Lloyd's Algorithm)* and the K-Means++ algorithm where the initialization overhead is directly proportional to the cluster count (k) for the data set. Our experimental results also serve the purpose of showing us that our initialization methodology has an over head very small compared to the total run time of the algorithm even for the smallest of chosen cluster counts. This combined with the proven choosing of better initialization centroids (as seen in Table 5.2 and Table 5.3 which show

Data-set Name	History Data-set cnt.	Cluster Count	Top 5 cnt	Hit Rate %age
Road Network	43	40	42	97.67%
		80	41	95.34%
		120	43	100%
		240	43	100%
Kegg Network	18	40	14	77.7%
		80	15	83.3%
		120	14	77.7%
		240	14	77.7%
US Gas Sensor Data	43	40	33	76.7%
		80	35	81.4
		120	35	81.4%
		240	34	79.1%
NotreDame	20	40	33	76.7%
		80	35	81.4
		120	35	81.4%
		240	34	79.1%
Tiny	20	80	20	100%
		120	18	90%
		240	16	80%
		360	17	85%
		480	18	90%
		600	18	90%
Uk Bench	20	80	20	100%
		120	18	90%
		240	16	80%
		360	17	85%
		480	18	90%
		600	18	90%
Caltec 101	20	80	20	100%
		120	18	90%
		240	16	80%
		360	17	85%
		480	18	90%
		600	18	90%

Table 5.1 Selection Hit Rate K-Means History Reuse

an overall improvement in both run time as well as the number of iterations taken for the data to converge) gives us a win win situation of a comprehensively better initialization methodology in

comparison to both the previously mentioned methods.

Referencing the above mentioned tables, we can also see that our selection mechanism takes only a very small percentage of the total run time for the algorithm for larger data sets ($< 0.1\%$) and while it may be a significant part of the initialization for smaller data sets the quality improvement in the initialization centroids is sufficient to offset the run time for the algorithm to be a viable initialization methodology for all data sets of varying sizes. Though speed up using our methodology is seen in all cases, it is best suited for larger data sets, as in that case our overhead for history best match selection almost becomes negligible.

Consistent Speedup for K-Means algorithm: The end requirement of any initialization algorithm is to ensure that we achieve a consistent speedup over previous approaches. While running our initialization algorithm for K-Means we were able to see consistent speedup in run-time of algorithm as compared to both K-Means and K-Means++ algorithm for all data-sets across a wide range of cluster counts ranging from 40 to 600. This is though dependent on quality of historical data-sets available. A bad match may lead to relatively bad starting points which in turn leads to an increase in time taken for convergence. Notwithstanding the above argument we were able to see consistently good performance by our algorithm for real data sets. This is largely due to the normalization of the data sets using the PCA based approach and then using cluster information based on this normalized data rather than the actual data. This also removes the need for the data to have Gaussian distribution. This gives a higher chance of finding a better historical match, thus giving consistent speedup over various data sets with different dimensionality (d) and different cardinality (n) for varying cluster counts k .

Figures 5.5 and 5.6 shows the range of speedup across various data sets with varying degrees of freedom as well as different dimensionality for the same cluster counts. We can clearly see from the graphs that our algorithm consistently out performs both the *random K-Means(Lloyd's Algorithm)* and the *K-Means++ algorithm* in run time performance.

Figures 5.3 and 5.4 depict the speedup seen by our algorithm over various cluster sizes for all data sets. We can see that the general trend is that speedup seen is in general higher for larger cluster count. This is particularly true because as the cluster count increases we see a speedup in both the initialization method as well as the actual run time because of the improvement seen in the initialization centroids.

Results seen in Table ?? and Table 5.2 together also show us that the speedups are consistently higher for larger cluster counts (k) in terms of both the run time of the algorithms and the number of iterations taken for the convergence of the algorithm to similar error rate. This is again an intuitive results which confirms our hypothesis: a consistent selection methodology independent of the cluster index size would lead to sizable improvement in run time. This is so because run time




Figure 5.3 Median Speedup, Small Data (k=240)




Figure 5.4 Median Speedup, Large Data (k=600)

Data-set	n	d	k	K-Means (Rand)	K-Means++	K-Means-HRU
Road Net.	1.0E4	3	40	79796	51830	12813
			80	116584	69990	16599
			120	149378	91757	20097
			240	231082	135479	154009
Kegg Net.	2.0E4	10	40	1.0534E5	0.7742E5	0.9984E5
			80	1.7604E6	1.4093E5	1.2473E5
			120	3.2291E5	2.4286E5	1.4068E5
			240	5.8874E5	3.2136E65	1.9502E5
US Gas	1.45E4	13	40	223657	135438	58243
			80	368228	243313	106190
			120	506618	281623	142108
			240	894857	524038	211386
NotreDame	1.45E4	13	40	2.1348E6	1.9956E6	1.3646E6
			80	3.874596E6	3.5660E6	2.4911E6
			120	5.480900E6	4.9390E6	3.2527E6
			240	1.0286953E6	6.7832E6	4.7613E6
Tiny	1.45E4	13	80	0.9610E7	0.9405E7	0.6631E7
			120	1.3710E7	1.3373E7	0.9516E7
			240	2.5315E7	2.4894E7	1.6639E7
			360	3.7154E7	3.4478E7	1.9563E7
			480	4.8749E7	3.7956E7	2.7585E7
			600	5.9941E7	4.5797E7	2.9498E7
Uk Bench	1.45E4	13	80	0.3512E8	0.3483E8	0.3347E8
			120	0.4976E8	0.4933E8	0.4727E8
			240	0.8277E8	0.8096E8	0.6589E8
			360	1.1910E8	1.0781E8	0.8302E8
			480	1.5552E8	1.3161E8	0.9558E8
			600	1.9259E8	1.4579E8	1.0992E8
Caltec 101	1.45E4	13	80	0.9231E7	1.0776E7	0.7587E7
			120	1.5687E7	1.5799E7	1.0574E7
			240	2.7573E7	2.7027E7	1.8120E7
			360	3.8998E7	3.7278E7	2.3259E7
			480	5.1231E7	4.9265E7	3.0193E7
			600	6.3268E7	5.4660E7	3.4443E7

Table 5.2 Runtime comparison: Kmeans, Kmeans++ and Kmeans-HRU (**All times in μ S*)

improvement is seen in both the initialization overhead as well as the actual run time of the algorithm. As the cluster count (k) increases, the amount of initialization overhead for both the above mentioned

Data-set	n	d	k	K-Means (Random)	K-Means++	K-Means-HRu
Road Net.	1.0E4	3	40	100	69	39
			80	100	52	27
			120	100	46	30
			240	100	79	82
US Gas	1.45E4	13	40	100	58	26
			80	100	59	32
			120	100	58	33
			240	100	55	37
NotreDame	1.45E4	13	40	100	78.5	69.5
			80	100	64	65
			120	100	93.5	60
			240	100	69	48.5
Tiny	1.45E4	13	80	100	100	100
			120	100	100	100
			240	100	100	87.5
			360	100	100	84
			480	100	87.5	74
			600	100	86.5	67
Uk Bench	1.45E4	13	80	100	100	85.5
			120	100	100	82
			240	100	100	75
			360	100	100	70.5
			480	100	79.5	63.5
			600	100	72	57.5
Caltec 101	1.45E4	13	80	100	100	84.5
			120	100	100	83.5
			240	100	100	74
			360	100	100	69
			480	100	90	61
			600	100	79	58.5

Table 5.3 Median Iteration count comparison for Kmeans, Kmeans++ and Kmeans with History Reuse.

initialization methodology becomes a significant part of the total run time for the algorithm.

Another interesting feature seen in the above box plots is the varying degrees of speedup seen for the various segments for same data-sets. Looking at the speedup that *K-Means with History Reuse* achieves over the original *K-Means (Lloyd's algorithm)* we can see that the speedup is generally spread out over a larger range as compared to the performance of the *K-Means++ algorithm* as

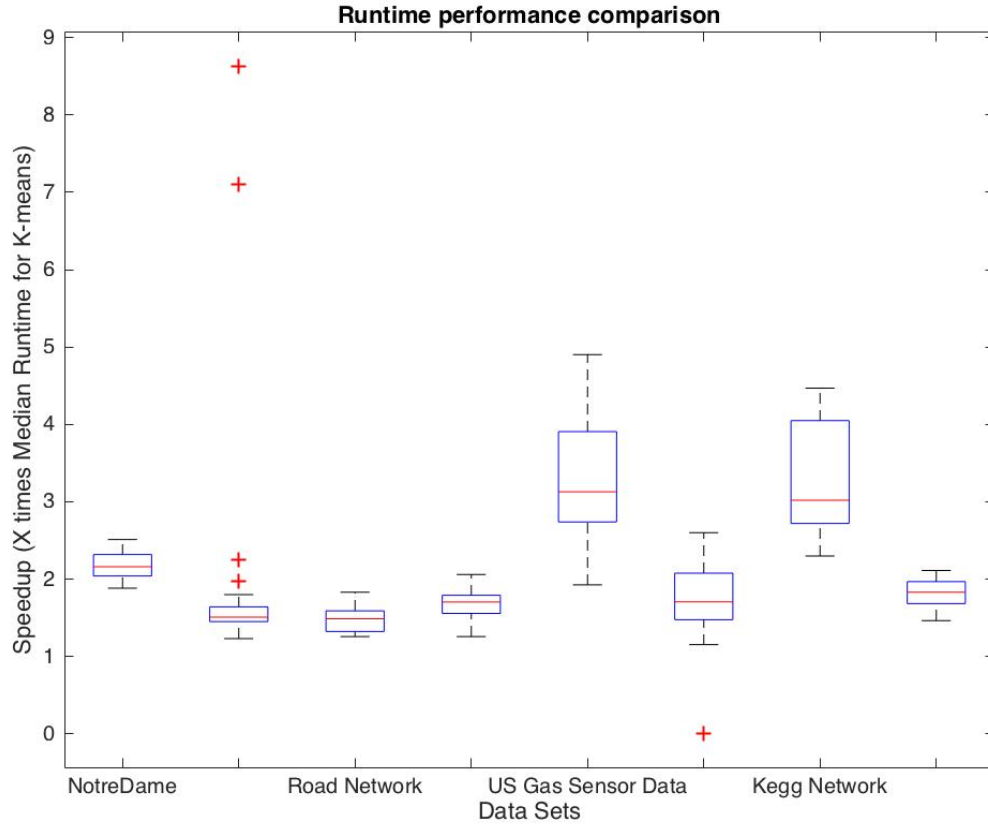


Figure 5.5 Median Speedup, Small Data (k=240)

compared to *llyod's Algorithm*. This thus, brings to fore our above made assumption about the quality of history data set available in terms of the probability of similarity. Our experiments prove our intuitive hypothesis that probabilistically similar data sets will tend to be classed together into similar clusters, thus making them the best history match data sets. The larger variance seen in the box plots above is testament to this. Normalization using *PCA* of data sets, tends to mitigate this to a large extent thus providing us the ability to use history data not along same planes as well and thus providing a consistent speedup, but a better match of normalized data would provide even higher speedups. We consistently saw that maximum speedups are seen when the probability of similarity is higher. For segment with high probabilistic similarity we saw speedups of up to and above **100X**.




Figure 5.6 Median Speedup, Large Data ($k=600$)

CHAPTER

6

RELATED AND FUTURE WORK

As we have seen in our experimentation the speedup in the run time for the algorithm is dependent of the quality (probability of similarity) of historically available data sets for selection. Better quality selection would lead to better speedups. Thus one future work could involve new methodology for maintaining best quality history data sets.

Another possibility is to explore this History Reuse methodology for other algorithms. The general use of determining best match history data set is generic in nature (i.e independent of the K-Means algorithm in itself) and thus can be extended to other algorithms such as the *SVD based SVM algorithm*. Only a minor change to the training step specific to K-Means is required in such a case. This thus can be used to prove the comprehensible usability of this algorithm for all algorithms in the same class.

On the systems side the History reuse methodology can be used for memory data validation in Non Volatile Memory (NVM). This would prevent multiple expensive re-writes to NVM from main memory in case of already existing data.

CHAPTER

7

CONCLUSION

In conclusion we would like to say that probabilistic metric for data-set comparison is one of the best metrics for this purpose. It tells us the exact degree to which we can expect the computation reuse to be successful for any historical data-set. It also represents a generic way of ranking data sets based on similarity and may have many more applications than just in history-reuse computation. We have also seen based on the above results that the above mentioned algorithm is handy for all data sets irrespective of the dimensionality, cardinality or cluster count for the data set. Though the algorithm shows speedups for all data sets for all cluster indexes we see a substantially higher speedup for larger cluster counts as this results in a speedup in both, the run time of the algorithm as well as the initialization step for the algorithm. The initialization is sped up considerably by our approach and this plays a vital part in the final speedup for larger cluster counts as initialization becomes a significant amount of the run time for both the *Random K-Means* and the *K-Means++* algorithm. Since this can be used as a drop in replacement for any initialization method, it can be used with any flavour of the original K-Means Algorithm. It preserves the semantic of the original K-means. These appealing properties, plus its simplicity, make it a practical replacement of the standard K-means as long as viable historical data sets are available for reuse.

BIBLIOGRAPHY

- [ER35] Einstein, A. & Rosen, N. "The particle problem in the general theory of relativity". *Physical Review* **48**.1 (1935), pp. 73–77.
- [GVL96] Golub, G. & Van Loan, C. *Matrix computations*. Johns Hopkins Univ Pr, 1996.
- [LK74] Larsen, E. & Keller, J. "Asymptotic solution of neutron transport problems for small mean free paths". *Journal of Mathematical Physics* **15** (1974), p. 75.

APPENDIX

APPENDIX

A

LOREM IPSUM

A.1 A First Section

A.1.0.0.1 Filler Text

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et

magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetur.

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.




Figure A.1 A figure in the appendix.

Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Donec odio elit, dictum in, hendrerit sit amet, egestas sed, leo. Praesent feugiat sapien aliquet odio. Integer vitae justo. Aliquam vestibulum fringilla lorem. Sed neque lectus, consectetur at, consectetur sed, eleifend ac, lectus. Nulla facilisi. Pellentesque eget lectus. Proin eu metus. Sed porttitor. In hac habitasse platea dictumst. Suspendisse eu lectus. Ut mi mi, lacinia sit amet, placerat et, mollis vitae, dui. Sed ante tellus, tristique ut, iaculis eu, malesuada ac, dui. Mauris nibh leo, facilisis non, adipiscing quis, ultrices a, dui.

Morbi luctus, wisi viverra faucibus pretium, nibh est placerat odio, nec commodo wisi enim eget quam. Quisque libero justo, consectetur a, feugiat vitae, porttitor eu, libero. Suspendisse sed mauris vitae elit sollicitudin malesuada. Maecenas ultricies eros sit amet ante. Ut venenatis velit. Maecenas sed mi eget dui varius euismod. Phasellus aliquet volutpat odio. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Pellentesque sit amet pede ac sem eleifend consectetur. Nullam elementum, urna vel imperdiet sodales, elit ipsum pharetra ligula, ac pretium ante justo a nulla. Curabitur tristique arcu eu metus. Vestibulum lectus. Proin mauris. Proin eu nunc eu urna hendrerit faucibus. Aliquam auctor, pede consequat laoreet varius, eros tellus scelerisque quam, pellentesque hendrerit ipsum dolor sed augue. Nulla nec lacus.

Suspendisse vitae elit. Aliquam arcu neque, ornare in, ullamcorper quis, commodo eu, libero. Fusce sagittis erat at erat tristique mollis. Maecenas sapien libero, molestie et, lobortis in, sodales eget, dui. Morbi ultrices rutrum lorem. Nam elementum ullamcorper leo. Morbi dui. Aliquam sagittis.

Table A.1 A table in the appendix.

System	Author
T _E X	Donald Knuth
L ^A T _E X	Leslie Lamport

Nunc placerat. Pellentesque tristique sodales est. Maecenas imperdiet lacinia velit. Cras non urna. Morbi eros pede, suscipit ac, varius vel, egestas non, eros. Praesent malesuada, diam id pretium elementum, eros sem dictum tortor, vel consectetur odio sem sed wisi.

A.2 A Second Section

Etiam ac leo a risus tristique nonummy. Donec dignissim tincidunt nulla. Vestibulum rhoncus molestie odio. Sed lobortis, justo et pretium lobortis, mauris turpis condimentum augue, nec ultricies nibh arcu pretium enim. Nunc purus neque, placerat id, imperdiet sed, pellentesque nec, nisl. Vestibulum imperdiet neque non sem accumsan laoreet. In hac habitasse platea dictumst. Etiam condimentum facilisis libero. Suspendisse in elit quis nisl aliquam dapibus. Pellentesque auctor sapien. Sed egestas sapien nec lectus. Pellentesque vel dui vel neque bibendum viverra. Aliquam porttitor nisl nec pede. Proin mattis libero vel turpis. Donec rutrum mauris et libero. Proin euismod porta felis. Nam lobortis, metus quis elementum commodo, nunc lectus elementum mauris, eget vulputate ligula tellus eu neque. Vivamus eu dolor.

Nulla in ipsum. Praesent eros nulla, congue vitae, euismod ut, commodo a, wisi. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Aenean nonummy magna non leo. Sed felis erat, ullamcorper in, dictum non, ultricies ut, lectus. Proin vel arcu a odio lobortis euismod. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Proin ut est. Aliquam odio. Pellentesque massa turpis, cursus eu, euismod nec, tempor congue, nulla. Duis viverra gravida mauris. Cras tincidunt. Curabitur eros ligula, varius ut, pulvinar in, cursus faucibus, augue.